



# Hindsight: A Prototype for a Flexible, Secure and Crash-safe Backup System

Johan Brinch    Morten Brøns-Pedersen  
Institute of Computer Science, University of Copenhagen



# Program



# Program

- Background



# Program

- Background
- Goals:
  - Flexible (front-end, back-end)
  - Efficient (time usage, space usage)
  - Secure (privacy, authentication)



# Program

- Background
- Goals:
  - Flexible (front-end, back-end)
  - Efficient (time usage, space usage)
  - Secure (privacy, authentication)
- System design:
  - Insert
  - Retrieve
  - Delete



# Program

- Background
- Goals:
  - Flexible (front-end, back-end)
  - Efficient (time usage, space usage)
  - Secure (privacy, authentication)
- System design:
  - Insert
  - Retrieve
  - Delete
- Demo!



# Program

- Background
- Goals:
  - Flexible (front-end, back-end)
  - Efficient (time usage, space usage)
  - Secure (privacy, authentication)
- System design:
  - Insert
  - Retrieve
  - Delete
- Demo!
- Questions?



# Background

Who are we, and what is this?





# Background

Who are we, and what is this?

- Students at DIKU



# Background

Who are we, and what is this?

- Students at DIKU
- Master thesis project



# Background

Who are we, and what is this?

- Students at DIKU
- Master thesis project
- Not just proof of concept



# Background

Who are we, and what is this?

- Students at DIKU
- Master thesis project
- Not just proof of concept
- Release candidate: August



# Background

Who are we, and what is this?

- Students at DIKU
- Master thesis project
- Not just proof of concept
- Release candidate: August (we hope)



# Goals

## Scope



# Goals

## Scope

- Personal computer



# Goals

## Scope

- Personal computer
- No sharing between systems





# Goals

## Scope

- Personal computer
- No sharing between systems
- Remote storage, over the Internet



# Goals

## Scope

- Personal computer
- No sharing between systems
- Remote storage, over the Internet
- Hostile environment



# Goals

## Scope

- Personal computer
- No sharing between systems
- Remote storage, over the Internet
- Hostile environment/user



# Goals

## Scope

- Personal computer
- No sharing between systems
- Remote storage, over the Internet
- Hostile environment/user
- (At least) daily backup



# Goals

A bunch of systems already, why a new one?



# Goals

A bunch of systems already, why a new one?

- Open source



# Goals

A bunch of systems already, why a new one?

- Open source
- Flexible back-end



# Goals

A bunch of systems already, why a new one?

- Open source
- Flexible back-end
- Global de-duplicity





# Goals

A bunch of systems already, why a new one?

- Open source
- Flexible back-end
- Global de-duplicity
- Scalability
  - Total log data
  - Per-snapshot overhead
  - Search and checkout



# Goals

A bunch of systems already, why a new one?

- Open source
- Flexible back-end
- Global de-duplicity
- Scalability
  - Total log data
  - Per-snapshot overhead
  - Search and checkout
- Deletion



# Goals

A bunch of systems already, why a new one?

- Open source
- Flexible back-end
- Global de-duplicity
- Scalability
  - Total log data
  - Per-snapshot overhead
  - Search and checkout
- Deletion
- Crash-safe



# System design

## Overview



# System design

## Overview

- A simple front-end API (key, metadata, value);



# System design

## Overview

- A simple front-end API (key, metadata, value);
- Idempotent key insertion and crash safety;



# System design

## Overview

- A simple front-end API (key, metadata, value);
- Idempotent key insertion and crash safety;
- Global de-duplication (content-based addressing);



# System design

## Overview

- A simple front-end API (key, metadata, value);
- Idempotent key insertion and crash safety;
- Global de-duplication (content-based addressing);
- Concurrent b-trees everywhere. No transactions;





# System design

## Overview

- A simple front-end API (key, metadata, value);
- Idempotent key insertion and crash safety;
- Global de-duplication (content-based addressing);
- Concurrent b-trees everywhere. No transactions;  
→ QuickCheck'ed... of course.




# System design

Hindsight system overview...




# System design

## Overview



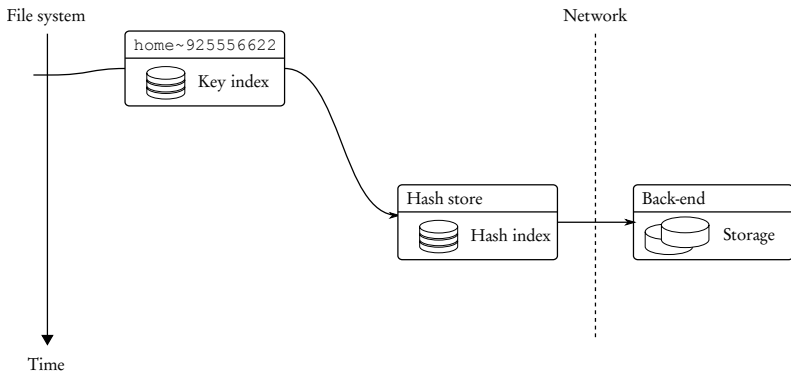
# System design

## Overview



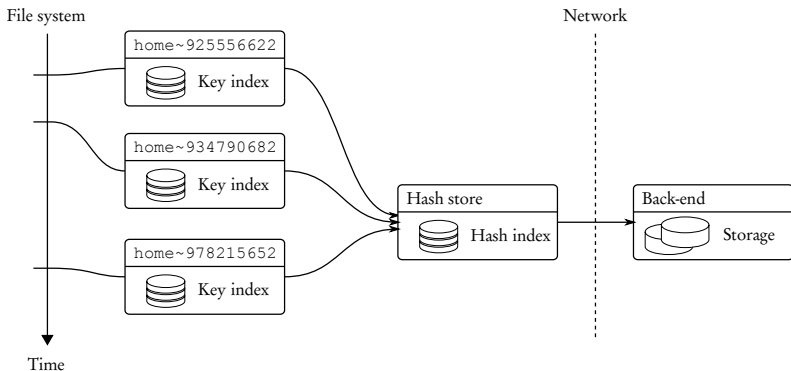
# System design

## Overview



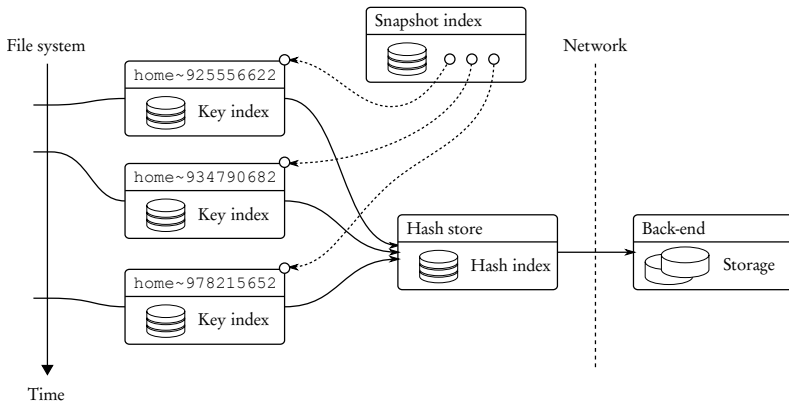
# System design

## Overview



# System design

## Overview



# System design


How we store files. . .






# System design

## Insertion



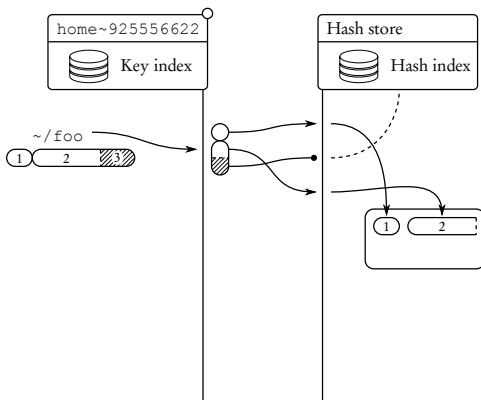
# System design

## Insertion



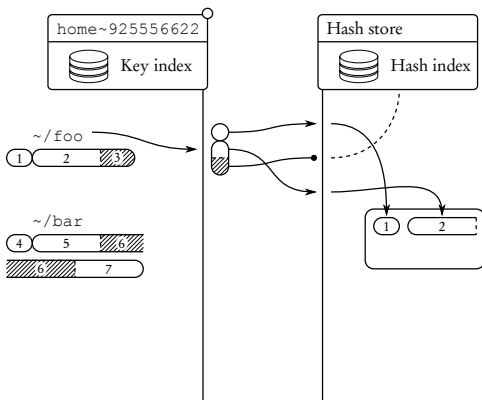
# System design

## Insertion



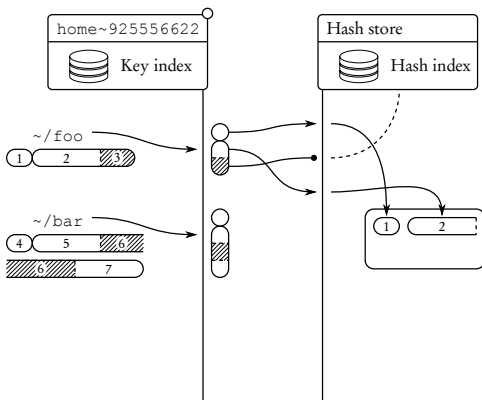
# System design

## Insertion



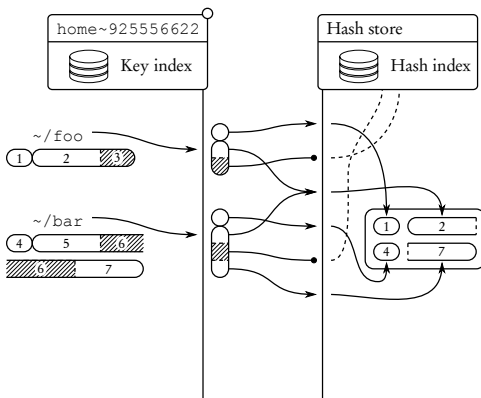
# System design

## Insertion



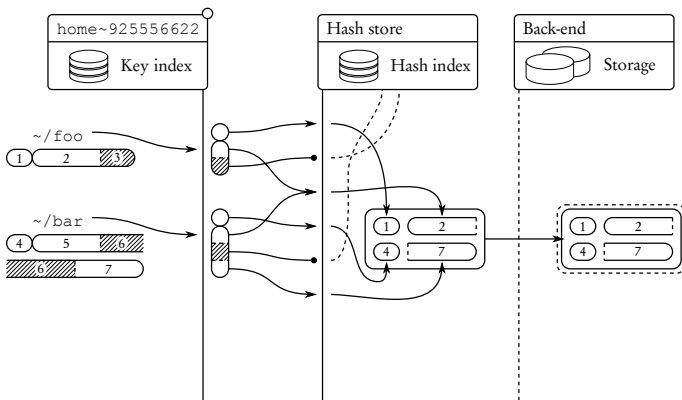
# System design

## Insertion



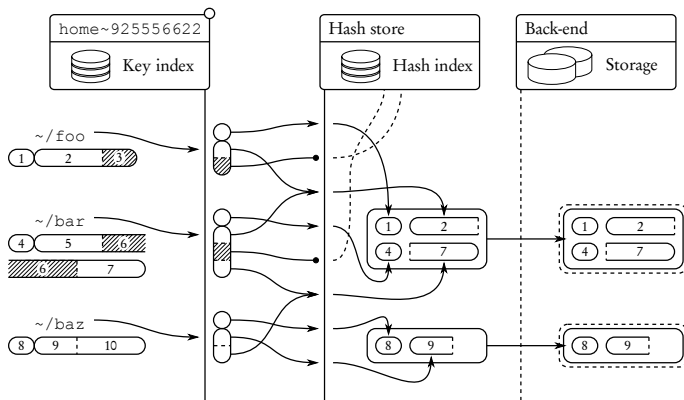
# System design

## Insertion



# System design

## Insertion





# System design

## Log data

- Now we just have the log data to take care of.



# System design

## Log data

- Now we just have the log data to take care of.
  - Much smaller than the data we backed up;



# System design

## Log data

- Now we just have the log data to take care of.
  - Much smaller than the data we backed up;
  - Multiple files with redundant data across snapshots;



# System design

## Log data

- Now we just have the log data to take care of.
  - Much smaller than the data we backed up;
  - Multiple files with redundant data across snapshots;
- Example: 10,278 files, 14 GB
  - key index: 64 files, 1.2 MB
  - hash index: 157 files, 1.3 MB



# System design

I wish we had some system for backing up files. . .





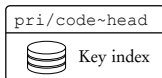
# System design



# System design


## Log data

Primary



# System design

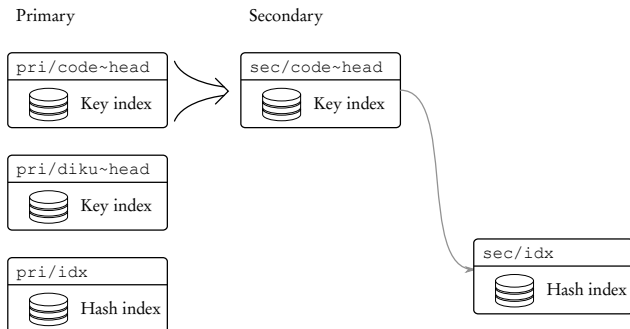
## Log data





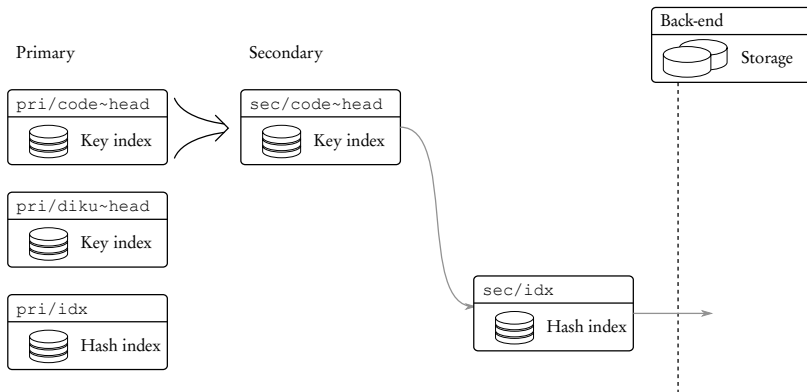
# System design

## Log data



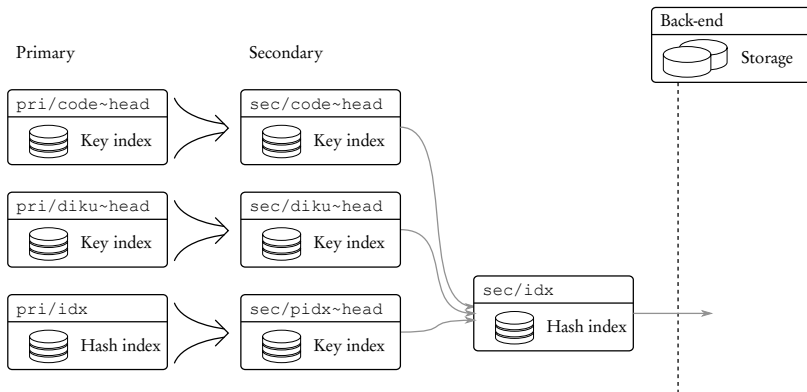
# System design

## Log data




# System design

## Log data




# System design

## Log data



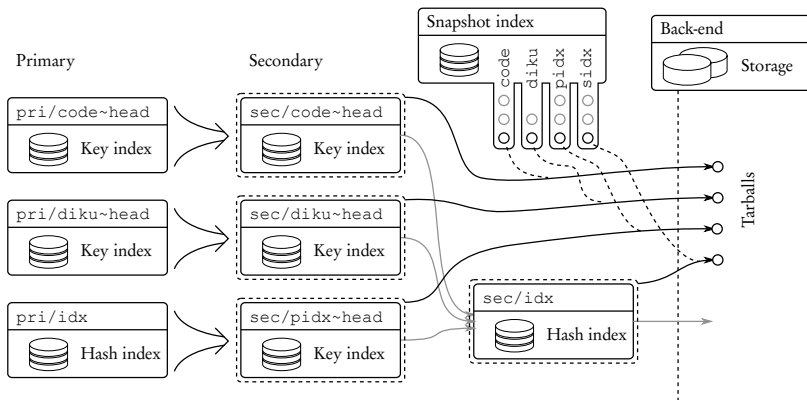
# System design

## Log data



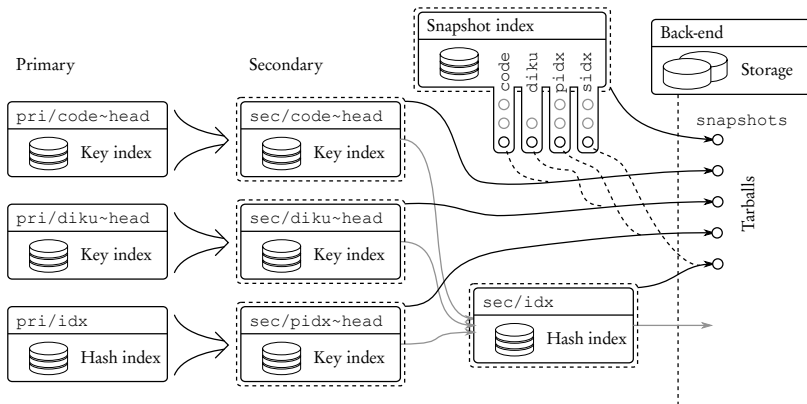
# System design

## Log data



# System design

## Log data



# System design

## Search and retrieve

- B-trees support lazy retrieval of nodes.





# System design

## Search and retrieve

- B-trees support lazy retrieval of nodes.
- Example using Amazon S3:



# System design

## Search and retrieve

- B-trees support lazy retrieval of nodes.
- Example using Amazon S3:

```
jos@laptux ~ > time hindsight list diku | wc -l  
104391  
real    26m00.28s
```



# System design

## Search and retrieve

- B-trees support lazy retrieval of nodes.
- Example using Amazon S3:

```
jos@laptux ~ > time hindsight list diku | wc -l  
104391  
real    26m00.28s
```

```
jos@laptux ~ > time hindsight list diku:speciale | wc -l  
14597  
real    3m20.21s
```



# System design

## Search and retrieve

- B-trees support lazy retrieval of nodes.
- Example using Amazon S3:

```
jos@laptux ~ > time hindsight list diku | wc -l  
104391  
real    26m00.28s
```

```
jos@laptux ~ > time hindsight list diku:speciale | wc -l  
14597  
real    3m20.21s
```

```
jos@laptux ~ > time hindsight listdir diku:speciale | wc -l  
10  
real    0m19.56s
```



# System design

## Deletion

- We want to delete unused file blocks.



# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.



# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.
  - But that is what the others are doing :/



# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.
  - But that is what the others are doing :/
- Reference sets! (aka reference listing)





# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.
  - But that is what the others are doing :/
- Reference sets! (aka reference listing)
  - Keep track of all references, not just the count.



# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.
  - But that is what the others are doing :/
- Reference sets! (aka reference listing)
  - Keep track of all references, not just the count.
  - Idempotent operations (insert / remove).



# System design

## Deletion

- We want to delete unused file blocks.
  - No transactions; no reference counting.
  - But that is what the others are doing :/
- Reference sets! (aka reference listing)
  - Keep track of all references, not just the count.
  - Idempotent operations (insert / remove).
  - Tests say the space overhead is acceptable.




# System design

Implementation:


Haskell alá Erlang: Processes and message passing.




# System design




# System design



# System design



# System design





# Demo!

Initiate kill-test now!



# Questions?

