



Hindsight: Flexible, secure and efficient backup

Johan Brinch and Morten Brøns-Pedersen

<http://github.com/mortenbp/hindsight>

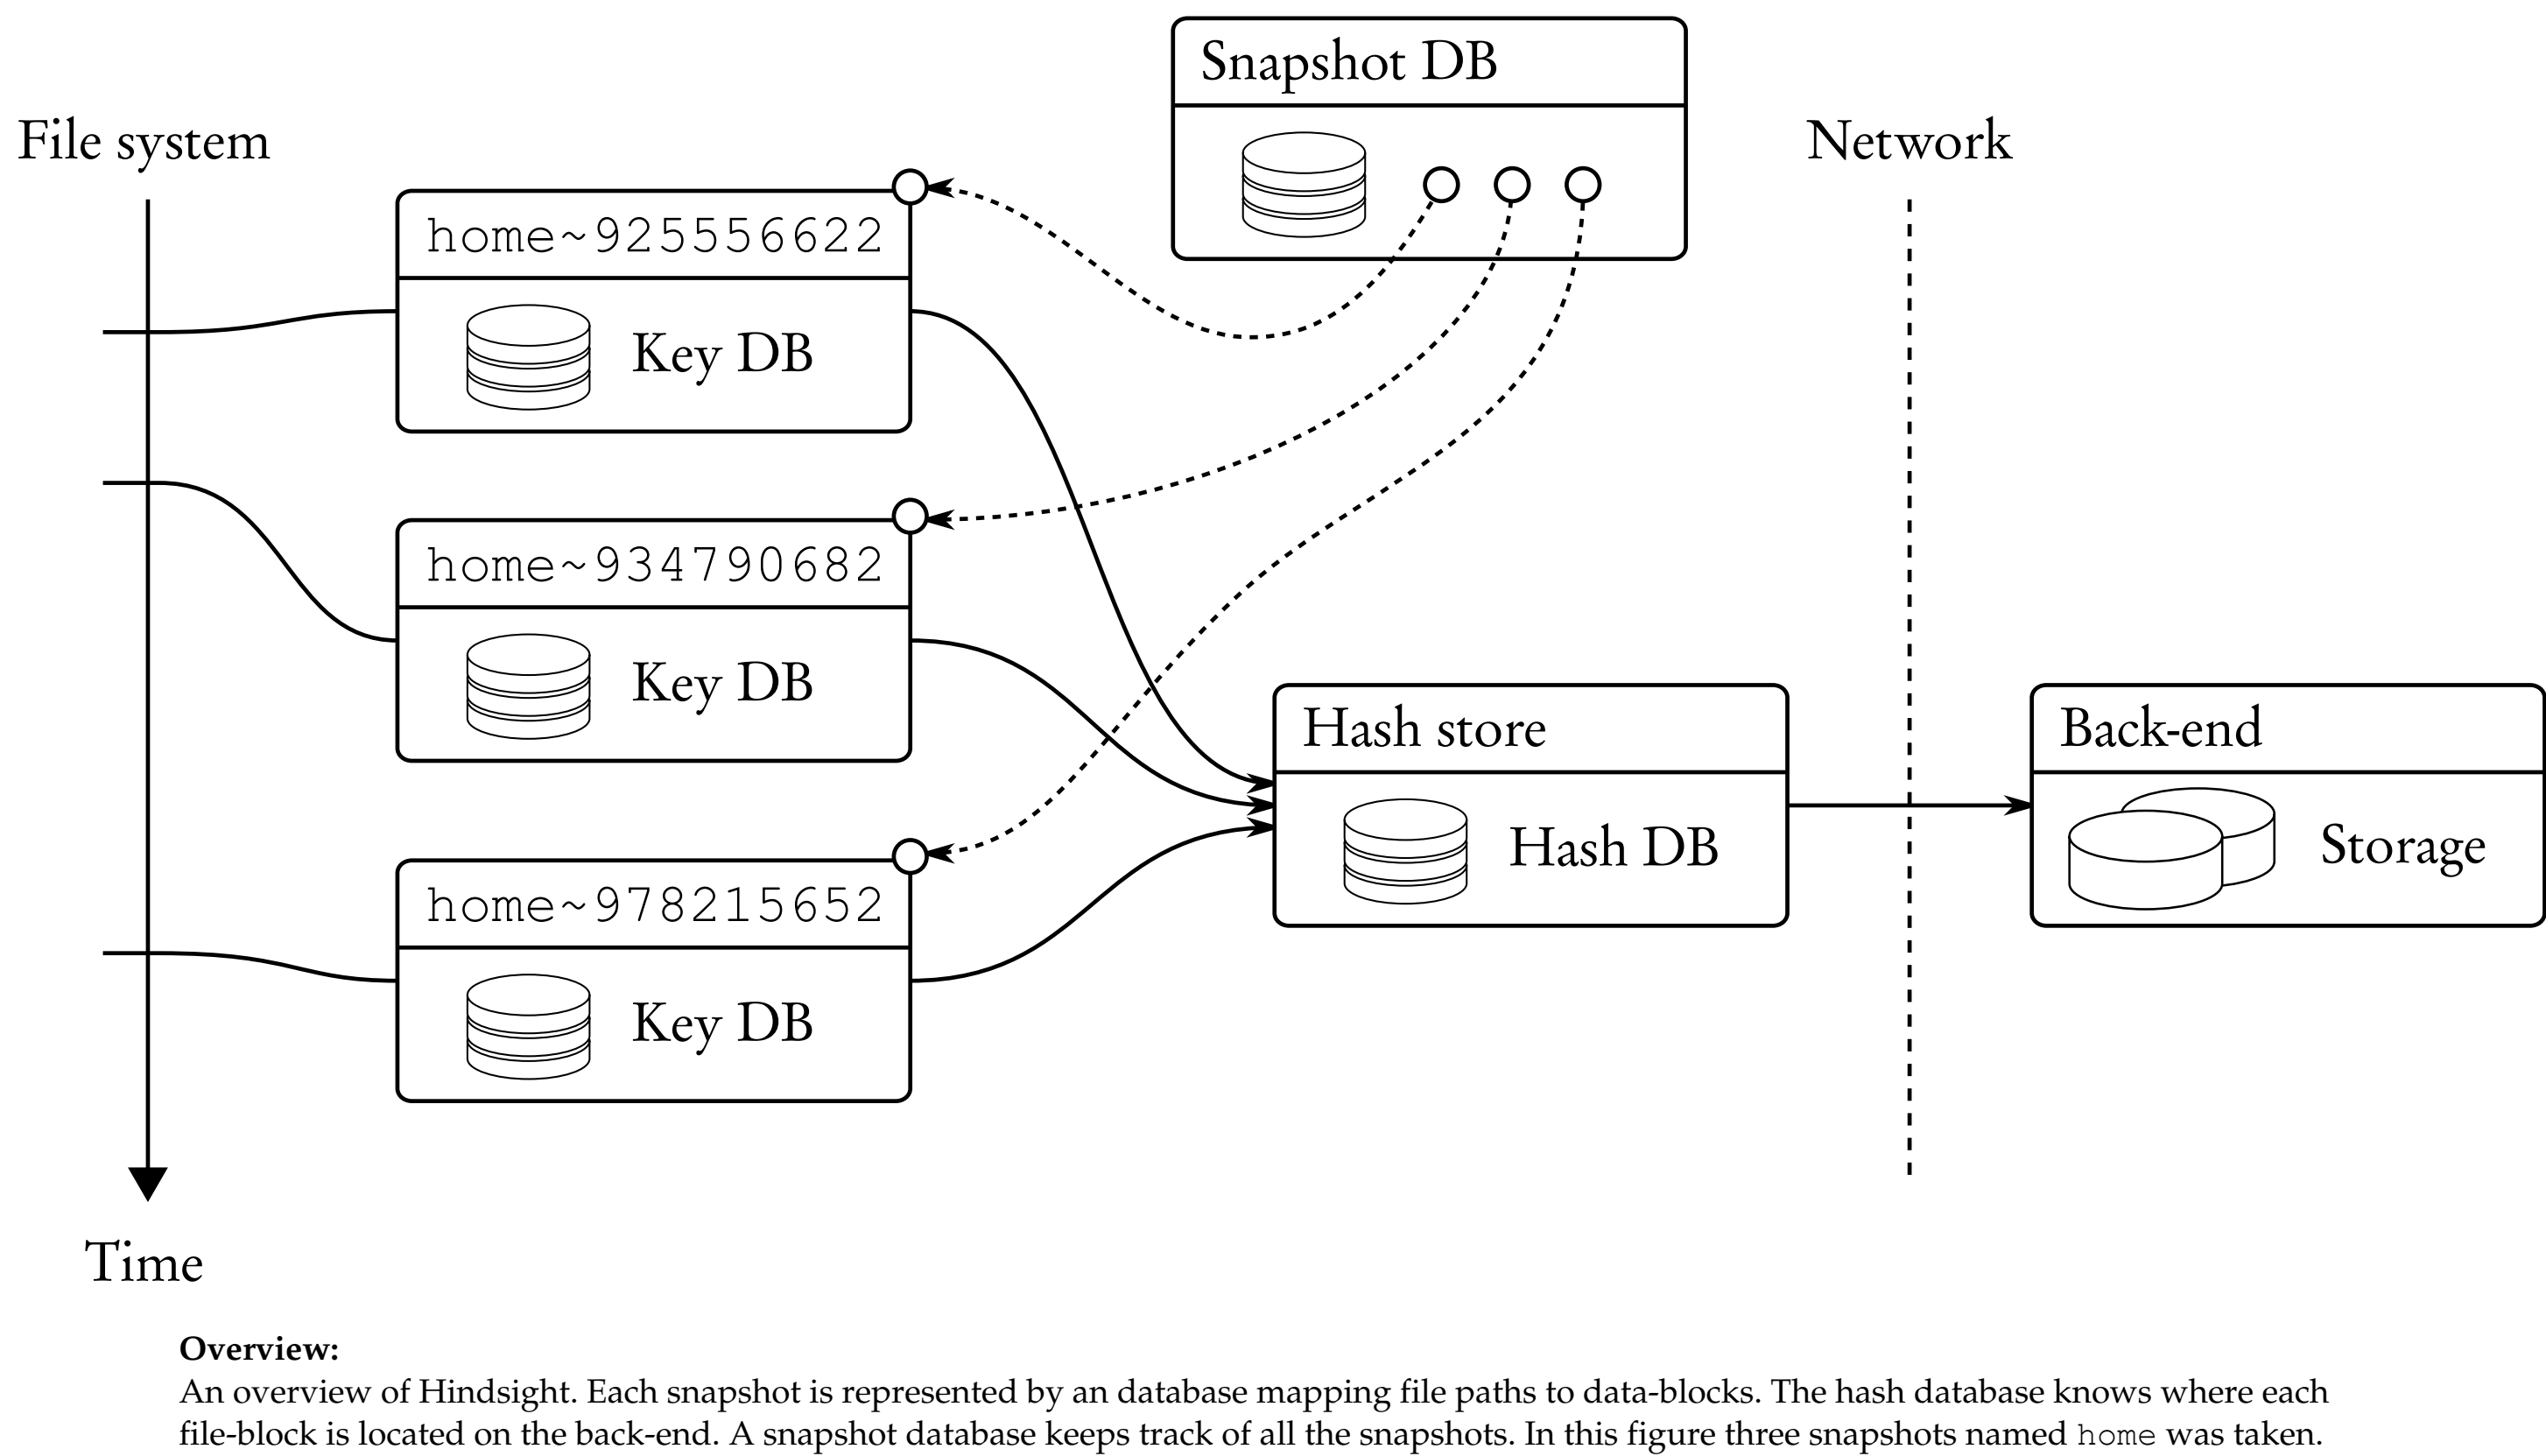
Abstract
Hindsight is a snapshot-based backup system providing guaranteed safe backup to untrusted sources such as Amazon S3 or DropBox. Our design features global de-duplicity through content-based indexing [1]; a means for exploiting data redundancy across all files in all snapshots.

Introduction
Hindsight is a backup system for filesystems. It backs up files and directories — and their metadata — to an external location. Multiple named snapshots are supported. A snapshot reflects the state of a directory at a particular time. Though usually only requiring a fraction of the space, each snapshot contains all files and directories that were in the file system when the snapshot was taken.

Hindsight is **flexible**. It can search through directory structures in snapshots and download just the directories or files you need; no more. This is especially useful when a range of snapshots must be searched for a file.


Hindsight is **efficient**. It uses compression to save space and eliminates redundant data across files; if more copies of the same file — or even just parts of it — exist, only the first one will travel to the back-end and be stored there. This is true across all snapshot ever taken. It processes several files concurrently to better utilise the disk and network, and it supports multi-core systems.

Hindsight is **secure**. It uses the NaCl [2] encryption library by Bernstein et al. to encrypt and authenticate the stored data. This prevents people who eavesdrop or has access to the back-end from reading or modifying the user's private data.



Crash safety
Hindsight uses concurrent b-trees [3] for all its internal state. This includes its two main databases: one maintaining files in the snapshots; and one maintaining known (i.e. previously saved) file-blocks.

For Hindsight to be crash-safe, these databases need to be crash-safe. We achieve this in a very simple, yet elegant way: due to the b-tree structure, we can simply write the trees button up: from leaf to root.



After a crash due to a failure (e.g. no more battery), a recovery mechanism sets in to recover as much of the prior state as possible. This allows Hindsight to resume a snapshot after a crash, and quickly skip to the stage at which it crashed or was forced to stop.

Hindsight can work on very hostile systems, that shuts down often. Tests show that Hindsight can work with just a few minutes between crashes, and still complete the snapshot eventually.

[1] Udi Manber. *Finding similar files in a large file system*. In Proceedings of the USENIX Winter 1994 Technical Conference on USENIX Winter 1994 Technical Conference, pages 2-2, Berkeley, CA, USA, 1994. USENIX Association.


[2] Daniel J. Bernstein, Tanja Lange, and Peter Schwabe. *The security impact of a new cryptographic library*. 2011.

[3] Kim S. Larsen and Rolf Fagerberg. *B-trees with relaxed balance*. In Proceedings of the 9th International Parallel Processing Symposium, pages 196-202. IEEE Computer Society Press, 1993.

[4] Amazon Web Services. *Amazon Simple Storage Service*, 2012. <http://goo.gl/wXJ1r>.

[5] J. Weizenbaum. *Knotted list structures*. Commun. ACM, 5:161-165, March 1962. ISSN 0001-0782. <http://goo.gl/PlmHU>.

[6] U. Maheshwari and B.H. Liskov. *Fault-tolerant distributed garbage collection in a client-server object-oriented database*. In Proceedings of the Third International Conference on Parallel and Distributed Information Systems, pages 239-248. IEEE, 1994.



Deletion
Global de-duplicity comes at a price. In order to share blocks between one file and another, we have to re-arrange the data and keep track of who's using what.

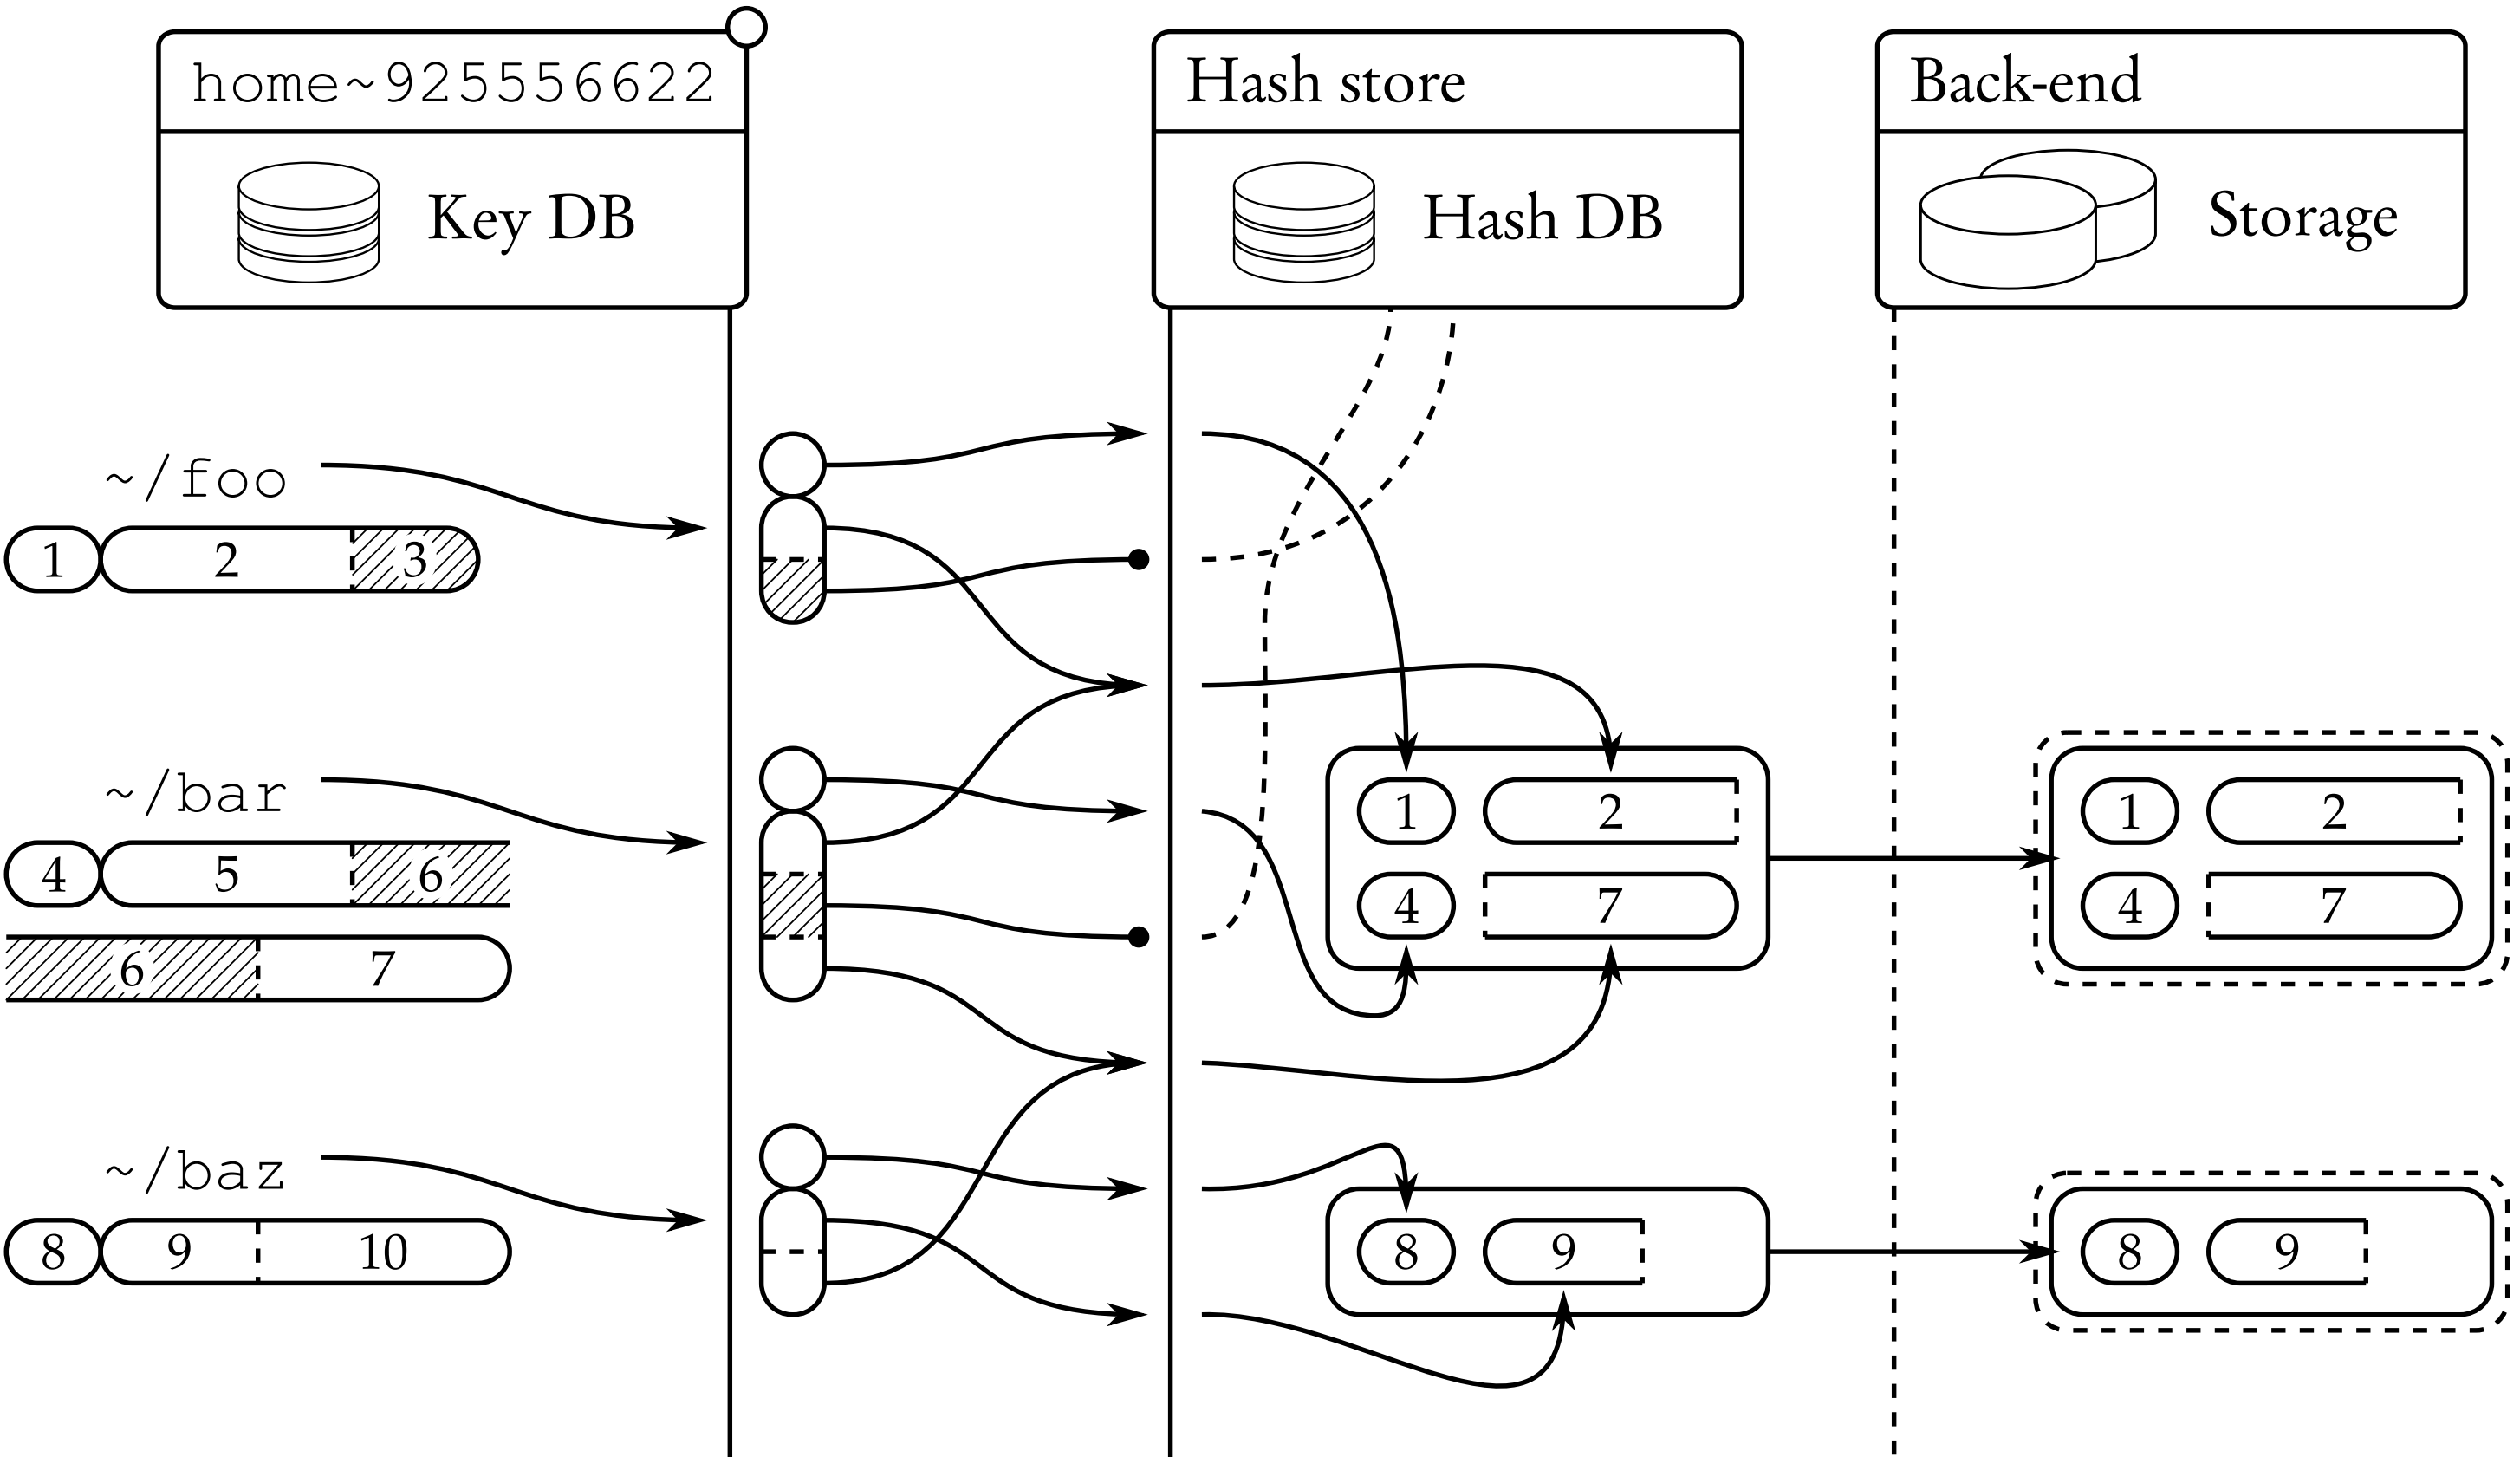
When a snapshot is deleted, we cannot simply delete all the blocks it is using. What if another snapshot is using them as well? Instead a garbage collector must analyse the block usage and remove the file-blocks that are referenced by no-one.

A classic approach for keeping track of references is called reference counting [5]. Here, each file-block is associated with a count of those referencing it. When the count goes to zero, the block can be safely deleted.

But reference counting requires transactions, which are incompatible with our b-trees. Instead, we solve the problem using reference sets [6]. Like reference counting, a reference set keeps track of references to a file-block. But instead of just maintaining the count, the set keeps track of everyone using the block. Thus adding or removing a reference several times is no longer harmful, which means transactions are no longer necessary.

The cost of maintaining reference sets — as opposed to reference counts — is more than made up for by a faster system, and a smaller space overhead. Our b-trees split naturally into multiple files, making it trivial to take advantage of the usually large amount of redundancy between b-trees belonging to different snapshots. This again means a faster system with a smaller per-snapshot space overhead.

Performance
The graph shows the amount of storage required to back up a virtual machine image (figure: space usage). The image is a KVM image of a machine running Ubuntu Linux 11.10. The first snapshot was taken right after the operating system finished installing. Then the system was updated, and software was installed and removed. A snapshot was performed after each such change.



The internals of Hindsight:
On the left are the files being backed up, and on the right is the back-end. The vertical lines represents mappings; the solid are databases stored locally as b-trees, and the dashed is the back-end API.